

# **AE 3340 Individual Project**

**Micaiah Smith-Pierce**  
903328232

*28 November, 2018*



# Contents

<b>1</b>	<b>Unconstrained Optimization</b>	<b>5</b>
1.1	Problems 1, 2, and 3 . . . . .	5
1.2	Problem 4 . . . . .	7
1.3	Problem 5 . . . . .	7
1.3.1	Part a . . . . .	7
1.3.2	Part b . . . . .	8
<b>2</b>	<b>Constrained Optimization</b>	<b>9</b>
2.1	Problem 6 . . . . .	9
2.2	Problem 7 . . . . .	9
2.3	Problem 8 . . . . .	11
2.4	Problem 9 . . . . .	11
2.4.1	Parts a - f . . . . .	11
2.4.2	Part g . . . . .	13
2.5	Problem 10 . . . . .	13



## Chapter 1

# Unconstrained Optimization

### 1.1 Problems 1, 2, and 3

The first three problems are analytic, so I did them by hand. They may be found on the following four pages.



## 1.2 Problem 4

Both Newton and Quasi-Newton methods essentially suppose that the objective function is quadratic, and then guess where the minimum is (or potentially which direction it is in). Although the assumption that the function is quadratic may seem arbitrary, in fact there is a good reason why it makes sense to assume that the function is quadratic, and nothing else. The reason is that the quadratic functions are the simplest ones which can actually have a minimum (without constraints). A linear function never has a strong minimum, and a cubic, or transcendental function would be unnecessarily complicated. So, an algorithm which can approximate the objective as a quadratic function is using the least possible amount of complexity and information in order to guess the location of the minimum. This is in contrast to the steepest descent method, which does not actually guess the location of the minimum, but simply attempts to decrease the objective as much as possible locally.

Moreover, both Newton and Quasi-Newton methods utilize the Hessian matrix, as well as the gradient. The Hessian matrix is the generalization of the second derivative of single-variable functions to multiple variables. Thus, just as the  $\nabla$  operator is the generalization of the derivative, the Hessian matrix  $H$  for a function  $f(\mathbf{x})$  can be defined in terms of  $\nabla$  and the vector outer product as:

$$H = \nabla \nabla^T f(\mathbf{x})$$

If, at a point  $\mathbf{x}_0$ ,  $f(\mathbf{x}_0)$ ,  $\nabla f(\mathbf{x}_0)$ , and  $H$  are known, and the function is assumed to be quadratic, then the location of the minimum,  $\mathbf{x}^*$  is known, and can be calculated as:

$$\mathbf{x}^* - \mathbf{x}_0 = -H^{-1} \nabla f(\mathbf{x}_0)$$

As stated in the MATLAB documentation. Here, the  $^{-1}$  denotes the matrix inverse. Realistically, however, the quadratic approximation may not be particularly accurate, so for a highly nonlinear function,  $\mathbf{x}^*$  may not be very close to the actual minimum unless  $\mathbf{x}_0$  is already close enough to the minimum that the higher-order derivatives do not significantly affect the behavior of the function. For that reason, rather than attempt to jump straight to the minimum, it may be advantageous to simply approximate the *direction* to the minimum, so that a standard line-search method can be used. In this case, the following similar equation gives the direction  $\mathbf{d}$  toward the minimum of the quadratic approximation:

$$\mathbf{d} = -H^{-1} \nabla f(\mathbf{x}_0)$$

The only difference between the Newton and Quasi-Newton methods is how the Hessian is computed (or estimated). The Newton method simply assumes that the Hessian is known, which means it must be computed either analytically, or using finite differences, i.e. sampling the function at points near  $\mathbf{x}_0$  in order to build a quadratic approximation. This can cost a lot of computation time if the function is expensive to evaluate, or the number of variables is large. For  $n$  variables,  $n^2$  points must be sampled to compute the Hessian by finite differences. The Quasi-Newton method solves this problem by using information from prior iterations of the algorithm. Rather than sample  $n^2$  new points at every iteration, it builds the quadratic approximation of  $f$  using the value of  $f$  at each previous point that the search has visited. With each iteration, it updates its previous estimate of the Hessian using the changes in the objective function's gradient. In this way, none of the information from previous iterations is wasted, which usually makes the algorithm more efficient.

## 1.3 Problem 5

### 1.3.1 Part a

The analytic gradient is clearly performs better than the estimated derivatives in all respect, and the steepest descent performs dramatically worse than either of the other two. The steepest descent algorithm only moves by extremely tiny increments, which is probably because the search direction does not take into account the curvature of the function. All of the search algorithms follow a curved path, but the steepest descent essentially must repeatedly stop to change directions, whereas the other two algorithms must do so less often, because the search direction anticipates the curvature of the path. So, the steepest descent algorithm ends up with more iterations and more function evaluations.

The estimated gradients method and analytic gradients method appear similar. However, the path that the estimated gradients takes utilizes shorter steps, and “zigzags” more than the analytic gradients. This is probably because the estimate of the gradient is not as accurate as the analytic computation, so

the search direction is subject to additional change, essentially because of uncertainty in the gradient computation. The estimated gradients method takes modestly more iterations to reach the solution, but it has nearly five times the number of function evaluations. This is probably because it must perform additional function evaluations during each iteration for the sole purpose of estimating the derivatives, which is completely unnecessary for the analytic gradient method.

### 1.3.2 Part b

Both the estimated and analytic gradient methods performed better in this case. This is probably because the starting point was located such that the algorithms did not have to perform as many iterations in the “valley”, which is the most challenging area. Yet, the most interesting thing about this example is that the analytic gradient took a completely different path from the estimated gradients function. The estimated gradients algorithm began searching straight down (toward the origin), whereas the analytic gradient took a path which is more in the positive  $x$ -direction. This is probably because the gradient vector was not perfectly vertical, which was detected and exploited by the analytic gradient algorithm, whereas the estimated gradient algorithm did not compute the gradient precisely enough to detect the asymmetry.



## Chapter 2

# Constrained Optimization

### 2.1 Problem 6

**Note:** Since computer code is inherently nondimensional, the problem is specified without units. In addition, let  $b$  and  $c$  be in feet.

**Minimize:**

$$f(b, c) = -\frac{L}{D}$$

**Subject to:**

$$g_1(b, c) = b - 10c \leq 0$$

$$g_2(b, c) = \frac{8000}{bc} - 40 \leq 0$$

$$g_3(b, c) = b - 55 \leq 0$$

### 2.2 Problem 7

Figure 2.1 below is my response to problem 7, parts a and b.

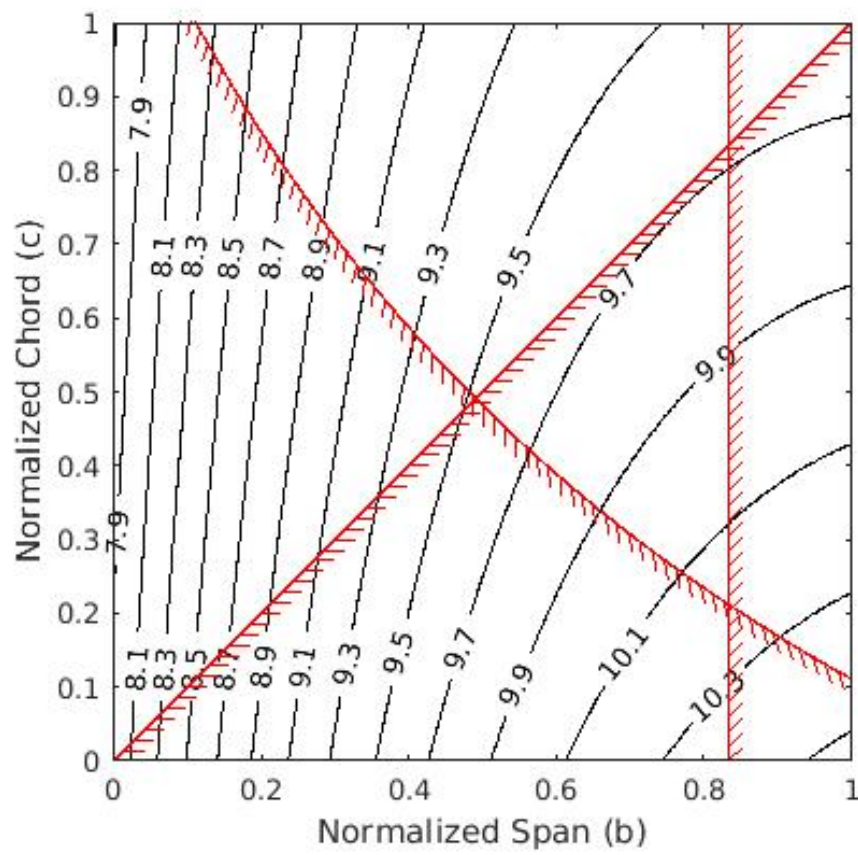


Figure 2.1: Answer to problem 7.

## 2.3 Problem 8

Algorithm	Iterations	Function Evaluations
interior-point	12	42
trust-region-reflective	Failed	Failed
sqp	7	24
sqp-legacy	7	24
active-set	4	15

The active-set algorithm was most efficient in terms of function evaluations. This agrees with the documentation, which seems to state that active-set is exceptionally fast in situations it is suited for.

## 2.4 Problem 9

### 2.4.1 Parts a - f

My response to parts a through f can be found on the following two handwritten pages.






Figure 2.2: The vectors  $\nabla f$  (black) and  $\lambda_1 \nabla g_1$  (red) at the computed minimum.

### 2.4.2 Part g

Yes, all of the KKT conditions are satisfied to an acceptable level of numerical precision. There are three KKT conditions. The first is that  $\mathbf{x}^*$  must be feasible. I clearly showed that this condition is met in part a. The second is that for any  $j$ ,  $\lambda_j g_j(\mathbf{x}^*) = 0$ . I showed this in part d. The final condition is that  $\nabla_{\mathbf{x}} \mathcal{L}(\mathbf{x}^*, \lambda) = \mathbf{0}$  (for the appropriate definition of  $\mathcal{L}$ ). I demonstrated this in part f. So, it has already been established that all three conditions are satisfied.

## 2.5 Problem 10

The vectors are plotted in Figure 2.2. It is apparent from this plot that  $\nabla \mathcal{L} = \mathbf{0}$  because the vector  $\lambda_1 \nabla g_1$  is equal in magnitude and opposite in direction from  $\nabla f$ . So,  $\nabla \mathcal{L}$ , which in this case is the sum of these two vectors, must be zero.